

Consider the Boston housing data set, from the MASS library. Consider column `medv` as a sample of the population prices of houses in the Boston area.

Assume that the data set is the parent sample. The population is the set of all houses in Boston.

- a) Use `t.test(d0$medv)` to find a 95% confidence interval for the mean of `medv`. Provide an estimate $\hat{\mu}$ for the population mean μ and of the standard deviation of that estimate $\hat{\mu}$
- b) Now estimate that standard deviation using the bootstrap.
- c) Based on your bootstrap estimate, find a 95% confidence interval for the mean of `medv`. Compare it to the results obtained using `t.test(d0$medv)`.
- d) Provide an estimate, $\hat{\mu}_{med}$, for the median of `medv` in the population.
- e) Estimate the standard error of $\hat{\mu}_{med}$. Unfortunately, there is no simple formula for computing the standard error of the median. Instead, estimate the standard error of the median using the bootstrap.
- f) Estimate the population tenth percentile $x_{0.1}$ of `medv`. Find the bootstrap estimate $\hat{x}_{0.1}$ and its standard error.

```

library(MASS)      # Boston data
library(boot)      # boot() boot.ci()

d0 = Boston
y = d0$medv

# classical estimate std error - MEAN
#=====
n = nrow(d0)        # 506
mean(y)             # [1] 22.53281
sd(y)/sqrt(n)       # [1] 0.4088611

aux=t.test(y,conf.level=0.95)
aux
#          One Sample t-test
# data:  y
# t = 55.111, df = 505, p-value < 2.2e-16
# alternative hypothesis: true mean is not equal to 0
# 95 percent confidence interval:
# 21.72953 23.33608
# sample estimates:
# mean of x
# 22.53281

t.test(y,conf.level=0.95)$conf
t.test(y,conf.level=0.95)$conf[1:2] # [1] 21.72953 23.33608

# Bootstrap estimate and bootstrap std error - MEAN
#=====

bfunction1 = function(data,index) mean(data[index]) # index are rows used to find mean()
bfunction1(y,1:n)      # 22.53281
bfunction1(y,1:n/2)    # 24.29703

# B=1000 bootstrap samples of size n=506
set.seed(1)
aux=boot(y, bfunction1, 1000)
aux
# ORDINARY NONPARAMETRIC BOOTSTRAP
# Bootstrap Statistics :
#   original      bias    std. error
# t1* 22.53281 0.008517589   0.4119374

str(aux)
# List of 11
# $ t0      : num 22.5
# $ t       : num [1:1000, 1] 22.5 22.3 22 22.6 23.1 ...

```

```

# $ R      : num 1000
# $ data   : num [1:506] 24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
# $ seed   : int [1:626] 403 624 -169270483 -442010614 -603558397 -222347416 1489374793 865871222 1
# $ statistic:function (data, index)
# ..- attr(*, "srcref")=Class 'srcref' atomic [1:8] 1 14 1 52 14 52 1 1
# .. ..- attr(*, "srcfile")=Classes 'srcfilecopy', 'srcfile' <environment: 0x1043b218>

# aux$t      the B resampling statistics
# mean(aux$t) the bootstrap is the mean of resampling stats (not reported)
# sd(aux$t)   sd of resampling stats (reported)

# resampling stats
head(aux$t)

# bootstrap mean, sd, bias

mean(aux$t)   # [1] 22.54132      # bootstrap mean (not reported)
sd(aux$t)     # [1] 0.4119374     # sd
mean(aux$t)-mean(y)      # bias
# [1] 0.008517589

# CI on mean
b1=boot(y, bfunction1, 1000)
boot.ci(b1,conf=0.95,type="basic")

# BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
# Based on 1000 bootstrap replicates
# Intervals :
# Level      Basic
# 95%        (21.67, 23.28)

```

```
# Bootstrap estimate and bootstrap std error - MEDIAN
#=====
median(y)          # [1] 21.2

bfunction2 = function(data,index) median(data[index])
bfunction2(y,1:n)   # [1] 21.2
bfunction2(y,1:n/2) # [1] 22.5

set.seed(1)
aux = boot(y,bfunction2,1000)
aux
# ORDINARY NONPARAMETRIC BOOTSTRAP
# Bootstrap Statistics :
#   original   bias    std. error
# t1*      21.2 -0.0025    0.374358

# bootstrap median
mean(aux$t)
# 21.1975
sd(aux$t)
# 0.374358
mean(aux$t)-median(y)  # bias
# -0.0025

# CI on median
b2=boot(y, bfunction2, 1000)
boot.ci(b2,conf=0.95,type="basic")
# BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
# Based on 1000 bootstrap replicates
# Intervals :
# Level      Basic
# 95%      (20.55, 21.95 )
```

```
# Bootstrap estimate and bootstrap std error - quantile
#=====
quantile(y,probs=c(0.1))    # 12.75

bfunction3 = function(data,index) quantile(data[index],probs=c(0.1))
bfunction3(y,1:n)
# 10%
#12.75
bfunction3(y,1:n/2)
# 10%
#15.6

set.seed(1)
aux=boot(y,bfunction3,1000)
aux
# ORDINARY NONPARAMETRIC BOOTSTRAP
# Bootstrap Statistics :
#   original    bias    std. error
#t1*      12.75  0.01005    0.505056

# bootstrap estimate
mean(aux$t)
# 12.76005
sd(aux$t)
# 0.505056
# bias
mean(aux$t)- quantile(y,probs=c(0.1))
# 10%
#0.01005

# CI on quantile
b3=boot(y,bfunction3, 1000)
boot.ci(b3,conf=0.95,type="basic")
# BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
# Based on 1000 bootstrap replicates
# Intervals :
# Level      Basic
# 95%      (12.05, 13.75 )
```